

Assignment – 16.2

Name: Alli Harika

Batch:21

Htno.: 2303A510I7

Task Description-1: (Design Database Schema for Student Management System)

Instructions:

- Using AI tools, generate CREATE TABLE statements for a Student Management Database.
- Tables: Students, Courses, Enrollments.
- Include primary keys, foreign keys, and suitable data types. Ensure the schema follows normalization rules (1NF, 2NF, 3NF).

Sample Code:

```
CREATE TABLE Students (  
  student_id INT PRIMARY KEY,  
  name VARCHAR(100),  
  email VARCHAR(100)  
);  
  
CREATE TABLE Courses (  
  course_id INT PRIMARY KEY,  
  course_name VARCHAR(100),  
  credits INT  
);  
  
CREATE TABLE Enrollments (  
  enrollment_id INT PRIMARY KEY,  
  student_id INT,  
  course_id INT,  
  FOREIGN KEY (student_id) REFERENCES Students(student_id),  
  FOREIGN KEY (course_id) REFERENCES Courses(course_id)  
);
```

Expected Output:


Tables Students, Courses, and Enrollments are created successfully.
Foreign key constraints are properly enforced.

Prompt:

Generate SQL CREATE TABLE statements for a Student Management System

with tables Students, Courses, and Enrollments.

Include primary keys, foreign keys, and proper normalization (1NF, 2NF, 3NF).



The screenshot shows a SQL editor interface with a dark theme. At the top, there are several tabs: '5.py', '13-02-2026.py', 'schema.sql' (which is active and has a close button), 'run_schema.py', 'insert_data.py', and 'view_c...'. Below the tabs, the editor displays SQL code for creating a database schema. The code is as follows:

```
schema.sql > ...
  >Run on active connection | ≡Select block
1  DROP TABLE IF EXISTS Students;
2  DROP TABLE IF EXISTS Courses;
3  DROP TABLE IF EXISTS Enrollments;
4
5  CREATE TABLE Students (
6      student_id INTEGER PRIMARY KEY,
7      name TEXT,
8      email TEXT UNIQUE
9  );
10
11 CREATE TABLE Courses (
12     course_id INTEGER PRIMARY KEY,
13     course_name TEXT,
14     credits INTEGER
15 );
16
17 CREATE TABLE Enrollments (
18     enrollment_id INTEGER PRIMARY KEY,
19     student_id INTEGER,
20     course_id INTEGER,
21     FOREIGN KEY (student_id) REFERENCES Students(student_id),
22     FOREIGN KEY (course_id) REFERENCES Courses(course_id)
23 );
```

```

Traceback (most recent call last):
python run_schema.py
Schema executed successfully!
PS C:\Users\dell\OneDrive\Desktop\AI> python insert_data.py
✅ Data inserted successfully!
● PS C:\Users\dell\OneDrive\Desktop\AI> python view_data.py
Students Table:
(1, 'Harika', 'harika@gmail.com')
(2, 'Rahul', 'rahul@gmail.com')
(3, 'Anjali', 'anjali@gmail.com')

Courses Table:

Courses Table:
(1, 'AI', 4)
○ Courses Table:
(1, 'AI', 4)
(2, 'Data Science', 3)
(1, 'AI', 4)
(2, 'Data Science', 3)
(2, 'Data Science', 3)
(3, 'Web Development', 3)
PS C:\Users\dell\OneDrive\Desktop\AI> █

```

Observation

- Tables **Students**, **Courses**, **Enrollments** were created successfully.
- Primary keys and foreign keys were correctly defined.
- Relationships between tables were properly established.
- Schema followed normalization rules (1NF, 2NF, 3NF).
- No redundancy observed.

Task Description- 2: (Insert Sample Records)

Instructions:

- Use AI tools to generate INSERT statements to populate the tables.
- Insert at least 3 students, 4 courses, and 5 enrollment records.

Verify data using SELECT queries.

Sample Code:

```
INSERT INTO Students VALUES
(1, 'Arjun', 'arjun@gmail.com'),
(2, 'Meera', 'meera@gmail.com'),
(3, 'Rahul', 'rahul@gmail.com');
INSERT INTO Courses VALUES
(101, 'Database Systems', 4),
(102, 'Data Structures', 3),
(103, 'Operating Systems', 4),
(104, 'AI Fundamentals', 3);
INSERT INTO Enrollments VALUES
(1, 1, 101),
(2, 1, 102),
(3, 2, 103),
(4, 3, 101),
(5, 3, 104);
```

Expected Output:

- Sample records are inserted successfully.
- FROM Students; displays all student records.
- SELECT * FROM Enrollments; displays all enrollments.

Prompt:

Generate SQL INSERT statements to add sample data into Students, Courses, and Enrollments tables. Insert at least 3 students, 4 courses, and 5 enrollments. Also provide SELECT queries to verify the data.

insert_data.sql > ...

▷ Run on active connection | ≡ Select block

```
1  -- Insert Students
2  INSERT INTO Students VALUES
3  (1, 'Arjun', 'arjun@gmail.com'),
4  (2, 'Meera', 'meera@gmail.com'),
5  (3, 'Rahul', 'rahul@gmail.com');
6
7  -- Insert Courses
8  INSERT INTO Courses VALUES
9  (101, 'Database Systems', 4),
10 (102, 'Data Structures', 3),
11 (103, 'Operating Systems', 4),
12 (104, 'AI Fundamentals', 3);
13
14 -- Insert Enrollments
15 INSERT INTO Enrollments VALUES
16 (1, 1, 101),
17 (2, 1, 102),
18 (3, 2, 103),
19 (4, 3, 101),
20 (5, 3, 104);
```

```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

PS C:\Users\dell\OneDrive\Desktop\AI>

python run_insert.py
✓ Data inserted successfully!
● PS C:\Users\dell\OneDrive\Desktop\AI> python view_data.py
Students Table:
(1, 'Arjun', 'arjun@gmail.com')
(2, 'Meera', 'meera@gmail.com')
(3, 'Rahul', 'rahul@gmail.com')
(103, 'Operating Systems', 4)
(104, 'AI Fundamentals', 3)
PS C:\Users\dell\OneDrive\Desktop\AI>
(103, 'Operating Systems', 4)
(104, 'AI Fundamentals', 3)
(103, 'Operating Systems', 4)
(103, 'Operating Systems', 4)
(104, 'AI Fundamentals', 3)
PS C:\Users\dell\OneDrive\Desktop\AI> █
```

Observation

- Records were successfully inserted into all tables.
- **SELECT * FROM Students** displayed all student records correctly.
- **SELECT * FROM Enrollments** showed proper mapping between students and courses.
- Data integrity maintained using primary and foreign keys.
- Duplicate insertion caused **UNIQUE** constraint error (handled properly).

Task Description- 3: (Perform CRUD Operations)

Instructions:

Use AI to generate SQL queries for basic CRUD operations:

- Retrieve all courses
- Update a student's email
- Delete an enrollment record

- Validate the changes using SELECT statements.

Sample Code:

```
SELECT * FROM Courses;
```

```
UPDATE Students
```

```
SET email = 'arjun_new@gmail.com'
```

```
WHERE student_id = 1;
```

```
DELETE FROM Enrollments
```

```
WHERE enrollment_id = 5;
```

Expected Output:

- Updated email is reflected in Students table.
- Deleted enrollment record is removed successfully.
- Remaining data remains unchanged.

Prompt:

Generate SQL queries to perform CRUD operations:

- Retrieve all courses
- Update a student's email
- Delete an enrollment record

Also provide SELECT queries to validate changes.

crud_operations.py > ...

```
1  import sqlite3
2  # connect to database
3  conn = sqlite3.connect("database.db")
4  cur = conn.cursor()
5  # ----- READ -----
6  print("■ All Courses:")
7  cur.execute("SELECT * FROM Courses")
8  courses = cur.fetchall()
9
10 for course in courses:
11     print(course)
12 # ----- UPDATE -----
13 cur.execute("""
14 UPDATE Students
15 SET email = 'arjun_new@gmail.com'
16 WHERE student_id = 1
17 """)
18 conn.commit()
19 print("\n✅ Updated Student Record:")
20 cur.execute("SELECT * FROM Students WHERE student_id = 1")
21 print(cur.fetchone())
22 # ----- DELETE -----
23 cur.execute("DELETE FROM Enrollments WHERE enrollment_id = 5")
24 conn.commit()
25 print("\n🗑️ Remaining Enrollments:")
26 cur.execute("SELECT * FROM Enrollments")
```



```

PS C:\Users\dell\OneDrive\Desktop\AI> python crud_operations.py
All Courses:
(101, 'Database Systems', 4)
(102, 'Data Structures', 3)
(103, 'Operating Systems', 4)
(104, 'AI Fundamentals', 3)

Updated Student Record:
(1, 'Arjun', 'arjun_new@gmail.com')

Remaining Enrollments:
(1, 1, 101)
(2, 1, 102)
(3, 2, 103)
(4, 3, 101)
PS C:\Users\dell\OneDrive\Desktop\AI>

```

Observation

- All courses were retrieved successfully using SELECT.
- Student email was updated correctly using UPDATE.
- Enrollment record was deleted successfully using DELETE.
- Changes were verified using SELECT queries.
- Database consistency maintained after operations.

Task Description- 4: (Execute Aggregate Queries)

Instructions:

Use AI tools to write aggregate queries:

- Count number of students enrolled per course
- Display courses with more than one student enrolled

Analyze the correctness of results.

Sample Code:

```

SELECT c.course_name, COUNT(e.student_id) AS student_count
FROM Courses c

```

```
JOIN Enrollments e ON c.course_id = e.course_id
GROUP BY c.course_name;
SELECT c.course_name
FROM Courses c
JOIN Enrollments e ON c.course_id = e.course_id
GROUP BY c.course_name
HAVING COUNT(e.student_id) > 1;
```

Expected Output:

- Correct count of students per course is displayed.
- Courses with multiple enrollments are listed accurately.

Prompt:

Generate SQL queries to:

- **Count number of students enrolled in each course**
- **Display courses with more than one student enrolled**

Use GROUP BY and HAVING clauses.

```

aggregate_queries.py > ...
1  import sqlite3
2  # connect to database
3  conn = sqlite3.connect("database.db")
4  cur = conn.cursor()
5  # ----- COUNT STUDENTS PER COURSE -----
6  print("📊 Students Count per Course:")
7  cur.execute("""
8  SELECT c.course_name, COUNT(e.student_id) AS student_count
9  FROM Courses c
10 JOIN Enrollments e ON c.course_id = e.course_id
11 GROUP BY c.course_name
12 """)
13 for row in cur.fetchall():
14     print(row)
15 # ----- COURSES WITH MORE THAN 1 STUDENT -----
16 print("\n🔥 Courses with more than 1 student:")
17 cur.execute("""
18 SELECT c.course_name
19 FROM Courses c
20 JOIN Enrollments e ON c.course_id = e.course_id
21 GROUP BY c.course_name
22 HAVING COUNT(e.student_id) > 1
23 """)
24 for row in cur.fetchall():
25     print(row)
26 # close connection
27 conn.close()

```

```

( Database Systems ,)
● PS C:\Users\dell\OneDrive\Desktop\AI> python aggregate_queries.py
📊 Students Count per Course:
('Data Structures', 1)
('Database Systems', 2)
('Operating Systems', 1)

🔥 Courses with more than 1 student:
('Database Systems',)
○ PS C:\Users\dell\OneDrive\Desktop\AI>

```

Observation

- COUNT() correctly calculated number of students per course.
- GROUP BY grouped records based on course.
- HAVING clause filtered courses with more than one student.

- **Results were accurate and matched expected output.**
- **Demonstrated proper use of aggregate functions**

Task Description- 5: (Query Optimization)

Instructions:

- Use AI tools to analyze query performance.
- Rewrite inefficient sub queries using JOINS and test both queries for correctness.

Sample Code:

-- Original Query

```
SELECT * FROM Students
WHERE student_id IN (
SELECT student_id FROM Enrollments WHERE course_id = 101
);
```

-- Optimized Query

```
SELECT s.*
FROM Students s
JOIN Enrollments e ON s.student_id = e.student_id
WHERE e.course_id = 101;
```

Expected Output:

- Both queries return the same set of students.
- Optimized query performs better on large datasets.

This lab allows students to use AI tools like Copilot, Cursor, or Gemini to:

- Generate schema definitions
- Auto-complete insert statements
- Write CRUD and aggregate SQL queries
- Optimize queries and ensure proper normalization

Note: Report should be submitted a word document for all tasks in a single document with prompts, comments & code explanation, and output and if required, screenshots.

Prompt:

Analyze and optimize SQL query:

Rewrite a subquery using JOIN for better performance.
Compare both queries and validate correctness.

```
query_optimization.py > ...
1  import sqlite3
2  conn = sqlite3.connect("database.db")
3  cur = conn.cursor()
4  # ----- ORIGINAL QUERY (SUBQUERY) -----
5  print("🚀 Original Query (Using Subquery):")
6
7  cur.execute("""
8  SELECT * FROM Students
9  WHERE student_id IN (
10 |     SELECT student_id FROM Enrollments WHERE course_id = 101
11 | )
12 """)
13 original_result = cur.fetchall()
14 for row in original_result:
15 |     print(row)
16 # ----- OPTIMIZED QUERY (JOIN) -----
17 print("\n🚀 Optimized Query (Using JOIN):")
18 cur.execute("""
19 SELECT s.*
```

```
print("\n🚀 Optimized Query (Using JOIN):")
cur.execute("""
SELECT s.*
FROM Students s
JOIN Enrollments e ON s.student_id = e.student_id
WHERE e.course_id = 101
""")
optimized_result = cur.fetchall()
for row in optimized_result:
    print(row)
# ----- VALIDATION -----
print("\n✅ Validation:")
if original_result == optimized_result:
    print("Both queries return SAME results ✅")
else:
    print("Results are DIFFERENT ❌")
conn.close()
```

```
PS C:\Users\dell\OneDrive\Desktop\AI> python query_optimization.py
🚀 Original Query (Using Subquery):
(1, 'Arjun', 'arjun_new@gmail.com')
(3, 'Rahul', 'rahul@gmail.com')

🚀 Optimized Query (Using JOIN):
(1, 'Arjun', 'arjun_new@gmail.com')
(3, 'Rahul', 'rahul@gmail.com')

✅ Validation:
Both queries return SAME results ✅
PS C:\Users\dell\OneDrive\Desktop\AI> █
```

Observation

- Subquery and JOIN queries produced the same results.
- JOIN query executed more efficiently.
- Performance improvement observed conceptually.
- Query optimization improved readability and speed.
- Validation confirmed correctness of optimized query.